

Evaluation-Led versus Governance-Led: Data Governance for Responsible AI in Asylum Decision-Making in the United Kingdom and Canada

July 2026

Abstract

Artificial intelligence (AI) is increasingly used to support asylum administration, where systems process highly sensitive personal information and inform high-risk decisions affecting fundamental rights. This study examines how data governance supports the responsible adoption of AI in asylum decision-making through a qualitative comparative case study of the United Kingdom and Canada, assessed against the seven characteristics of trustworthy AI in the NIST AI Risk Management Framework. The analysis argues that the two jurisdictions embody distinct governance models. The United Kingdom pursues an evaluation-led approach, publicly evaluating deployed tools such as Asylum Case Summarisation (ACS) and Facial Age Estimation (FAE). Canada pursues a governance-led approach centred on the Directive on Automated Decision-Making and mandatory Algorithmic Impact Assessments. Evidence from the UK evaluation indicates that AI-assisted tools can reduce processing times and improve administrative efficiency — but the analysis argues that such gains are defensible only within a strict governance framework with meaningful human oversight. Each model manages a different failure mode: evaluation-led governance surfaces errors after deployment but makes them visible and correctable, whereas governance-led control reduces risk before deployment but produces little public evidence of operational performance. The study concludes that trustworthy AI in asylum systems requires elements of both models — lifecycle data governance combined with published operational evaluation — and that AI should augment rather than replace human decision-making in refugee status determination.

1. Introduction

Artificial intelligence has transitioned from a technological innovation to a foundational infrastructure of modern public administration. Governments increasingly adopt AI to modernise services, improve operational efficiency, and support evidence-based decision-making across sectors including healthcare, finance, and immigration.^{1,2} Beyond administrative decision-making, AI is being applied across the broader refugee response: the World Bank, for example, has developed an AI-powered forecasting model that

¹ Stanford Human-Centered Artificial Intelligence (2026) The AI Index 2026 Annual Report. Stanford, CA. Available at: <https://hai.stanford.edu/ai-index/2026-ai-index-report>

² UNDP (2025) Human Development Report 2025 — A matter of choice: people and possibilities in the age of AI. New York: UNDP. Available at: <https://hdr.undp.org/content/human-development-report-2025>

analyses socio-economic and conflict-related indicators to predict refugee movements before large-scale displacement occurs.³

The adoption of AI within asylum systems is driven not only by innovation but by operational pressure. Across the European Union, approximately 1.20 million asylum applications were awaiting a first-instance or final decision at the end of February 2026, with Germany alone reporting more than 291,000 pending applications.⁴ Asylum procedures are complex, lengthy, and resource-intensive, and pressure from rising applications and accumulating backlogs adds to this challenge — leaving applicants waiting for decisions that determine their legal status and access to rights and services.⁵ Against this backdrop, governments are exploring AI-assisted technologies to improve efficiency and consistency, while the European Parliament⁶ and the European Commission⁷ stress that deployment must remain transparent, accountable, subject to meaningful human oversight, and compliant with fundamental rights.

Deployment in this domain raises acute governance challenges. AI systems in the asylum administration process handle highly sensitive personal information and support decisions that can have potentially life-changing consequences. Asylum and migration systems are explicitly classified as high-risk under Annex III(7) of the European Union Artificial Intelligence Act, precisely because they affect individuals in particularly vulnerable situations.⁸ Effective data governance is therefore the foundation on which the quality, transparency, accountability, and trustworthiness of AI-assisted decision-making depend.

This study examines that foundation through a comparative analysis of the United Kingdom and Canada and pursues two research questions. **RQ1:** How do the United Kingdom and Canada govern data across the AI lifecycle in asylum decision-making? **RQ2:** What are the trade-offs between the United Kingdom's evaluation-led approach and Canada's governance-led approach, and what do they imply for other jurisdictions adopting AI in high-risk public services?

2. Methodology

This study adopts a qualitative comparative case study methodology, enabling an in-depth examination of governance practices within their real-world contexts and the identification of similarities and

³ World Bank (2025) AI-powered refugee forecasting: preparing for refugee movements before they happen. Available at: <https://www.worldbank.org/en/topic/fragilityconflictviolence/brief/ai-powered-refugee-forecasting-preparing-for-refugee-movements-before-they-happen>

⁴ Eurostat (2026) Asylum applications in February 2026. News article, 18 May 2026. Available at: <https://ec.europa.eu/eurostat/web/products-eurostat-news/w/ddn-20260518-1>

⁵ European Parliament (2025) Artificial intelligence in asylum procedures in the EU. Briefing PE 775.861, by C. Dumbrava. Brussels: EPRS. Available at: [https://www.europarl.europa.eu/thinktank/en/document/EPRS_BRI\(2025\)775861](https://www.europarl.europa.eu/thinktank/en/document/EPRS_BRI(2025)775861)

⁶ European Parliament (2025) — see note 5.

⁷ European Commission (2026) European asylum and migration management strategy. Communication, 29 January 2026. Available at: https://home-affairs.ec.europa.eu/european-asylum-and-migration-management-strategy_en

⁸ European Union (2024) Regulation (EU) 2024/1689 laying down harmonised rules on artificial intelligence (Artificial Intelligence Act). Official Journal of the European Union, L series. Available at: <https://eur-lex.europa.eu/eli/reg/2024/1689/oj>

differences between jurisdictions. The analysis draws on publicly available legislation, government policies, official guidance, evaluation reports, and academic literature. It does not assess the technical performance or predictive accuracy of AI models; instead, it examines the governance arrangements that support their development, deployment, and oversight.

The United Kingdom and Canada were selected because both have adopted AI-assisted technologies in asylum and immigration decision-making while implementing distinct governance approaches. The UK Home Office has publicly evaluated specific deployed applications — Asylum Policy Search (APS), Asylum Case Summarisation (ACS), and Facial Age Estimation (FAE) — whereas Immigration, Refugees and Citizenship Canada (IRCC) has adopted a broader governance model centred on the Directive on Automated Decision-Making (DADM) and mandatory Algorithmic Impact Assessments (AIAs). Several other jurisdictions also deploy AI in asylum or migration management — Germany's Federal Office for Migration and Refugees, for example, uses automated language and dialect identification to support origin assessment — but the United Kingdom and Canada offer two advantages for comparative documentary analysis. First, both publish sufficient official documentation, including evaluations and impact assessments, to support analysis from public sources. Second, both share comparable common-law, Westminster-style administrative traditions, which help isolate the governance model itself — rather than differences in legal culture — as the principal point of contrast.

Findings are interpreted against the legal frameworks applicable in each jurisdiction — the UK General Data Protection Regulation (UK GDPR) and the Canadian federal framework, including the DADM and privacy legislation — with the EU AI Act used as an internationally recognised benchmark for high-risk AI governance, and the UNHCR approach to artificial intelligence providing a complementary humanitarian perspective. The NIST AI Risk Management Framework (AI RMF 1.0) was selected as the analytical lens, rather than the EU AI Act itself, because it is technology-neutral, jurisdiction-neutral, and structured around characteristics of trustworthiness rather than the compliance obligations of any single legal system — allowing both jurisdictions to be assessed on equal terms.⁹

3. Analytical Framework and Comparative Overview

The NIST AI RMF identifies seven interrelated characteristics of trustworthy AI: valid and reliable; safe; secure and resilient; accountable and transparent; explainable and interpretable; privacy-enhanced; and fair with harmful bias managed.¹⁰ Table 1 summarises how each jurisdiction addresses these characteristics and what public evidence exists of implementation. The table reveals the pattern developed throughout this paper: the United Kingdom generates public evidence primarily through

⁹ National Institute of Standards and Technology (2023) Artificial Intelligence Risk Management Framework (AI RMF 1.0). NIST AI 100-1. Gaithersburg, MD: NIST. Available at: <https://nvlpubs.nist.gov/nistpubs/ai/NIST.AI.100-1.pdf>

¹⁰ National Institute of Standards and Technology (2023) — see note 9.

operational evaluation of deployed tools, while Canada generates it through mandatory pre-deployment governance documentation.

Table 1. Comparative overview of UK and Canadian governance mechanisms across the NIST AI RMF characteristics.

NIST characteristic	United Kingdom	Canada	Public evidence of implementation
Valid and reliable	AI-ready datasets guidance (2026); UK GDPR Recital 71 (accuracy in automated processing)	Directive on Automated Decision-Making (DADM) quality-assurance requirements	UK: published ACS evaluation identifying errors. Canada: none publicly available
Safe	Five cross-sectoral principles; FAE tested (not yet deployed), must not be sole basis for age decisions	Algorithmic Impact Assessments (AIAs) scaled to impact level	UK: FAE guidance and testing; external HRW critique. Canada: published AIAs
Secure and resilient	AI Cyber Security Code of Practice; secure-by-design guidance	DADM controls; exploration of Canadian data residency	Neither publicly verifiable beyond policy documents
Accountable and transparent	ICO accountability requirements; AI Playbook for government	DADM responsibility assignment; proposed AIDA (lapsed 2025)	UK: published tool evaluations. Canada: published AIAs
Explainable and interpretable	Explainability principle in 2023 regulatory framework	DADM explanation requirements scaled to impact level	Limited in both jurisdictions
Privacy-enhanced	UK GDPR principles; ICO guidance on AI and data protection	Privacy Act; IRCC AI Strategy privacy commitments	Documentary evidence only in both
Fair, with harmful bias managed	Equality Act 2010; Public Sector Equality Duty; ICO bias guidance	AIA bias assessment; IRCC responsible-AI principles	UK: ACS evaluation. Canada: AIA disclosures

4. Comparative Analysis

4.1 Reliability and Safety

Validity, reliability, and safety are treated together because in asylum decision-making they converge on the same question: can AI-generated outputs be depended upon in decisions with potentially life-changing consequences? AI systems should produce accurate, consistent outputs for their intended purpose and minimise foreseeable harm throughout their lifecycle,¹¹ which in turn requires high-quality, representative, well-governed data at every stage.

¹¹ National Institute of Standards and Technology (2023) — see note 9.

The United Kingdom addresses these requirements through a combination of data governance guidance and operational evaluation. Government guidance on AI-ready datasets requires data to be documented, representative, and fit for purpose;¹² the AI regulatory framework identifies safety, security and robustness among its five cross-sectoral principles;¹³ and UK GDPR Recital 71 requires appropriate mathematical and statistical procedures to minimise error in automated processing. Critically, the Home Office has published evaluations of its deployed tools. The ACS evaluation found that although the system improved administrative efficiency, some AI-generated summaries contained inaccuracies and omitted information, with varying levels of user confidence.¹⁴ FAE, by contrast, has been tested but is not yet in operational use in the immigration system; the Home Office's guidance, issued ahead of deployment, states that machine-learning age estimates must not be the sole basis for determining an individual's age.¹⁵ Human Rights Watch¹⁶ has nevertheless argued that errors in AI-assisted age estimation could expose asylum-seeking children to significant harm if they are wrongly classified as adults, losing access to child-specific protections.

Canada addresses the same requirements primarily before deployment. The DADM requires federal institutions to conduct AIAs, implement quality-assurance measures, and monitor automated systems in proportion to their risk level.¹⁷ The IRCC Artificial Intelligence Strategy commits to systems that are equitable, explainable, privacy-protecting, and technically sound, and confirms that automated tools cannot refuse applications — refusal decisions remain with immigration officers.¹⁸

Trade-off verdict. The UK's evaluation-led approach discovered real reliability failures — but only after deployment, and the published evaluation is what made them visible and correctable. Canada's governance-led approach is designed to prevent such failures upfront, yet no published operational evaluation exists against which its effectiveness can be tested. The UK model accepts visible, correctable error; the Canadian model accepts unverifiable prevention. Neither alone satisfies Articles 9, 10, and 15

¹² UK Government (2026) Guidelines and best practices for making government datasets ready for AI. GDS/DSIT, January 2026. Available at: <https://www.gov.uk/government/publications/making-government-datasets-ready-for-ai/guidelines-and-best-practices-for-making-government-datasets-ready-for-ai>

¹³ UK Government (2023) A pro-innovation approach to AI regulation. White Paper. Available at: <https://www.gov.uk/government/publications/ai-regulation-a-pro-innovation-approach>

¹⁴ Home Office (2025) Evaluation of AI trials in the asylum decision making process. Research and analysis, April 2025. Available at: <https://www.gov.uk/government/publications/evaluation-of-ai-trials-in-the-asylum-decision-making-process>

¹⁵ Home Office (2026) Facial age estimation: using AI to support initial age decisions (a guide). Guidance, May 2026. Available at: <https://www.gov.uk/government/publications/facial-age-estimation>

¹⁶ Human Rights Watch (2026) The UK's discriminatory AI experiment on child refugees. 19 June 2026. Available at: <https://www.hrw.org/news/2026/06/19/the-uks-discriminatory-ai-experiment-on-child-refugees>

¹⁷ Treasury Board of Canada Secretariat (2023) Directive on Automated Decision-Making. Available at: <https://www.tbs-sct.canada.ca/pol/doc-eng.aspx?id=32592>

¹⁸ Immigration, Refugees and Citizenship Canada (2026) IRCC's Artificial Intelligence Strategy. Available at: <https://www.canada.ca/en/immigration-refugees-citizenship/corporate/transparency/artificial-intelligence-strategy.html>

of the EU AI Act, which require continuous risk management, data governance, and accuracy across the whole lifecycle.¹⁹

4.2 Security and Privacy

Security, resilience, and privacy are assessed together because both jurisdictions address them through the same instruments, and because in both cases the evidence available is documentary rather than operational. AI-assisted asylum systems process identity records, biometric data, family relationships, health information, and accounts of persecution; maintaining confidentiality, integrity, and availability — and processing no more personal data than necessary — is essential to protecting applicants and preserving public trust.²⁰ Article 10 of the EU AI Act requires appropriate data governance for training, validation, and testing datasets, and Article 15 requires resilience against errors, faults, and unauthorised interference.²¹

The United Kingdom's framework combines the UK GDPR principles — lawfulness, purpose limitation, data minimisation, storage limitation, integrity, and accountability — with the ICO's guidance that organisations remain responsible for protecting personal data throughout the AI lifecycle,²² and with the AI Cyber Security Code of Practice, which recommends securing training data, models, infrastructure, access controls, and incident response.²³ Canada's framework combines the Privacy Act, the DADM's proportionate governance and monitoring requirements, and the IRCC AI Strategy, which identifies cybersecurity, privacy protection, and responsible information management as central commitments and explores Canadian data residency for AI-enabled services.^{24,25} The Government of Canada's Guide on the Use of Generative AI adds requirements to protect sensitive information and document AI use.²⁶

Trade-off verdict. On paper, the jurisdictions are near parity: the UK leans towards secure-by-design technical guidance, Canada towards institutional information management. But for both, the limited availability of technical documentation or independent audit prevents any assessment of how these controls operate in practice. In this domain, the two models converge on the same weakness — security and privacy assurances that cannot be externally verified. From a data governance perspective,

¹⁹ European Union (2024) — see note 8.

²⁰ National Institute of Standards and Technology (2023) — see note 9.

²¹ European Union (2024) — see note 8.

²² Information Commissioner's Office (2023) Guidance on AI and data protection. Available at: <https://ico.org.uk/for-organisations/uk-gdpr-guidance-and-resources/artificial-intelligence/guidance-on-ai-and-data-protection/>

²³ UK Government (2025b) AI Cyber Security Code of Practice. Available at: <https://www.gov.uk/government/publications/ai-cyber-security-code-of-practice>

²⁴ Treasury Board of Canada Secretariat (2023) — see note 17.

²⁵ Immigration, Refugees and Citizenship Canada (2026) — see note 18.

²⁶ Government of Canada (2024) Guide on the use of generative artificial intelligence. Available at: <https://www.canada.ca/en/government/system/digital-government/digital-government-innovations/responsible-use-ai/guide-use-generative-ai.html>

meaningful assurance would additionally require role-based access restricted to authorised personnel and documented staff training on handling sensitive asylum data.

4.3 Accountability and Explainability

Accountability and explainability are treated together because explainability is what makes accountability operational: a caseworker who cannot interpret an AI output cannot meaningfully take responsibility for a decision that relies on it, and an applicant who cannot obtain an explanation cannot meaningfully challenge one. Responsibility for AI-assisted decisions must remain with identifiable organisations and authorised decision-makers, not with systems.²⁷ Articles 13 and 14 of the EU AI Act require high-risk systems to operate with transparency and meaningful human oversight, and Article 86 strengthens the right of affected individuals to receive explanations of decisions informed by high-risk AI.²⁸

In the United Kingdom, the ICO requires organisations to remain accountable for AI-supported decisions and to provide meaningful information about processing, while the Artificial Intelligence Playbook for government requires clear governance responsibilities, documented use of AI, and human oversight.²⁹ The ACS evaluation gives these principles operational content: because some summaries omitted relevant information, responsibility for reviewing outputs demonstrably remains with caseworkers and the Home Office — accountability cannot be transferred to the system or its vendor when errors occur.³⁰

Canada assigns accountability through formal governance: the DADM requires institutions to assign responsibility for automated systems, conduct AIAs, and provide explanations scaled to the system's impact level, with the AIA process documenting design, intended use, risks, and oversight arrangements.³¹ IRCC confirms that officers, not systems, make refusal decisions.³² Canada's proposed Artificial Intelligence and Data Act (AIDA) would have imposed additional obligations on high-impact systems, but the bill lapsed when Parliament was prorogued in 2025 and had not been reintroduced at the time of writing.³³

Trade-off verdict. Both jurisdictions locate accountability in the public authority. The difference lies in what each makes visible: the UK publishes what its tools got wrong, which is accountability after the fact; Canada publishes what its institutions assessed and assigned, which is accountability before the fact. For an asylum applicant, however, neither model yet guarantees the full chain that transparency requires —

²⁷ National Institute of Standards and Technology (2023) — see note 9.

²⁸ European Union (2024) — see note 8.

²⁹ UK Government (2025a) Artificial Intelligence Playbook for the UK Government. Available at: <https://www.gov.uk/government/publications/ai-playbook-for-the-uk-government>

³⁰ Home Office (2025) — see note 14.

³¹ Treasury Board of Canada Secretariat (2023) — see note 17.

³² Immigration, Refugees and Citizenship Canada (2026) — see note 18.

³³ Innovation, Science and Economic Development Canada (2023) The Artificial Intelligence and Data Act (AIDA) — companion document. Available at: <https://ised-isde.canada.ca/site/innovation-better-canada/en/artificial-intelligence-and-data-act-aida-companion-document>

knowing that AI was used, understanding how it processed their information, and having an accessible route to human review or redress.

4.4 Fairness and Harmful Bias

Fairness warrants standalone treatment because it is where the limits of both governance models are most exposed. Every asylum claim is unique and assessed within its own legal, humanitarian, and factual context: applicants differ in country of origin, ethnicity, religion, language, gender, personal experience, and supporting evidence. Constructing training datasets that adequately represent the full diversity of refugee experience is therefore both technically and ethically challenging — and because refugee movements, patterns of persecution, and country conditions continually evolve, datasets that were representative at deployment degrade over time. Article 10 of the EU AI Act accordingly requires training, validation, and testing data to be relevant, sufficiently representative, and designed to identify and mitigate bias.³⁴

The United Kingdom addresses fairness through the Equality Act 2010 and the Public Sector Equality Duty, which require public authorities to consider impacts on protected groups; through fairness as a cross-sectoral principle of the AI regulatory framework;³⁵ and through ICO guidance that requires bias assessment of datasets and AI-assisted processing.³⁶ Canada addresses it through the AIA process, which requires bias-related risks to be assessed and mitigated in proportion to impact, and through the IRCC Strategy's commitment to equitable systems.^{37,38}

Trade-off verdict. Neither jurisdiction publishes evidence that training data for asylum-facing systems is representative, regularly reviewed, or updated as displacement patterns change — the UK's evaluations examined output quality, not dataset composition, and Canada's AIAs disclose assessed risk, not data lineage. Fairness in this domain, therefore, currently rests on assertion rather than demonstration in both models. Continuous dataset review, periodic bias re-assessment, and meaningful human oversight are the minimum conditions under which either approach could substantiate its fairness claims.

5. Discussion: What the Two Models Teach Each Other

The comparative analysis supports a central claim: the United Kingdom and Canada have not simply implemented the same principles differently — they have adopted governance models that manage different failure modes, and each model's strength is the other's weakness.

³⁴ European Union (2024) — see note 8.

³⁵ UK Government (2023) — see note 13.

³⁶ Information Commissioner's Office (2023) — see note 22.

³⁷ Treasury Board of Canada Secretariat (2023) — see note 17.

³⁸ Immigration, Refugees and Citizenship Canada (2026) — see note 18.

The evaluation-led model makes errors visible. The ACS errors are public knowledge precisely because the Home Office evaluated and published; that visibility is what enables correction, external scrutiny, and institutional learning. But it also reveals a deeper concern than inaccuracy alone: omission. When an AI summarisation tool silently removes information from an asylum interview, the caseworker has no way of knowing what was removed or on what basis the system judged it dispensable. A caseworker cannot challenge what they cannot see. Unless users understand how the system selects and discards information, human oversight risks becoming procedural rather than meaningful — a signature on an output no one is equipped to question.

The governance-led model prevents risk upfront but resists external verification. Canada's AIAs and the DADM's proportionate controls are, on paper, the more systematic architecture. Yet because no operational evaluations of IRCC's asylum-facing systems are public, the outputs of these systems are unknown outside the institution. This is not only a transparency concern — it raises the question of how accountability would function if the system were producing flawed outputs that no published mechanism exists to surface. It is also, candidly, a limitation of the present study: the Canadian half of this comparison rests on what Canada says its governance does, not on what its systems do.

Field experience adds a perspective that documents alone cannot provide. The author has worked for over a decade in large-scale refugee registration and verification operations. In real asylum settings, data is difficult from the moment it is created. Many refugees and asylum seekers arrive with no documentation at all, so dates of birth and other personal details, such as place of birth, are often estimated rather than recorded. Many flee alone and are joined by family members later in the country of asylum, so family records change over time and require continuous updating. Names are also spelt differently across systems, and duplicate records accumulate with repeated registrations. When agencies run large verification exercises, they routinely find gaps between what the database says and the situation on the ground. An AI system trained on this data inherits all of these problems. Policy documents describe how data should be managed, but they rarely record where the data came from, how reliable it actually is, or the conditions under which it was collected. Credible data governance for AI in asylum systems must therefore go below the policy layer and into the operational layer — because that is where the quality of the data, and of every AI output built on it, is really decided.

For a jurisdiction deploying AI in asylum administration today, the practical implication is to combine the models rather than choose between them. From the governance-led model: mandatory pre-deployment impact assessment, documented responsibility, and controls proportionate to risk. From the evaluation-led model: published operational evaluation, with testing and error recording in every trial. And beneath both: training data that is diverse and representative of the population the system will affect, with transparency across the full AI lifecycle — documented processes from data collection through deployment, monitoring, and review. Prevention without published evaluation cannot demonstrate that it worked; evaluation without lifecycle governance discovers failures only after people have been affected by them.

Finally, the UNHCR approach to artificial intelligence adds a dimension that neither government framework fully captures: the humanitarian do-no-harm principle.³⁹ Government frameworks govern AI for lawful and efficient administration — they ask whether a system complies with data protection law, whether risks have been assessed, and whether oversight is in place. Do-no-harm asks a different question: whether the intervention itself, even when lawful and well-governed, could harm the person it processes. In asylum contexts this distinction matters because errors that are administratively minor can be existentially serious for the individual — a misclassified age, an omitted detail in an interview summary, or a mishandled record can affect protection outcomes for someone with no alternative recourse. A protection-centred governance model would therefore treat the asylum seeker not merely as a data subject with legal rights, but as a person whose safety is the ultimate measure of whether an AI deployment has succeeded.

6. Limitations and Future Research

This study is based on publicly available legislation, policies, official guidance, evaluation reports, and academic literature. Its findings therefore assess documented governance frameworks, not the technical systems themselves, and several limitations follow. First, public documentation does not identify the datasets used to train, validate, or test the systems examined, so their representativeness, quality, and update frequency could not be assessed. Second, operational AI outputs were not available, preventing independent evaluation of consistency, fairness, or reliability across case types — a limitation that weighs most heavily on the Canadian case, as discussed above. Third, internal architecture, data lineage, audit logs, standard operating procedures, and retention practices were not accessible. Fourth, the study includes no empirical data from caseworkers or asylum applicants, so the gap between documented governance and day-to-day implementation could not be measured.

These limitations define the future research agenda. Semi-structured interviews with Home Office and IRCC staff could establish how governance frameworks operate in practice and whether staff are equipped to interpret, evaluate, and challenge AI-generated outputs. Access to internal governance documentation — data lineage, model documentation, audit logs, and dataset review records — would show how data governance principles are implemented across the lifecycle. Research should also examine organisational oversight structures: whether dedicated bodies monitor performance, review model updates, investigate errors, and manage complaints. Finally, and most importantly, future work should incorporate the perspectives of asylum seekers themselves: whether they understand when AI supports their case, whether they receive meaningful information about how their data is processed, and whether mechanisms for clarification, human review, and redress are genuinely accessible.

³⁹ UNHCR (2024) UNHCR's approach to artificial intelligence. Available at: <https://www.unhcr.org/digitalstrategy/ai-approach/>

7. Conclusion

This study asked how data governance supports the responsible adoption of AI in asylum decision-making (RQ1) and what trade-offs distinguish the UK and Canadian approaches (RQ2). On the first question, the analysis across all seven NIST AI RMF characteristics shows that trustworthy AI in asylum systems rests on data governance across the full lifecycle: representative and continuously reviewed datasets, documented processes, monitoring, privacy protection, bias management, and meaningful human oversight. On the second, the two jurisdictions embody complementary but individually insufficient models: the UK's evaluation-led approach makes failures visible and correctable at the cost of discovering them in deployment, while Canada's governance-led approach builds systematic preventive control at the cost of public verifiability. A jurisdiction adopting AI in high-risk public services should combine mandatory pre-deployment assessment with published operational evaluation, resting on transparent, diverse, lifecycle-governed data.

AI can materially improve asylum administration — supporting document analysis, interview summarisation, translation, triage, and the prioritisation of vulnerable applicants. The evidence reviewed suggests that AI-assisted tools can help reduce the persistent backlogs facing asylum authorities, and these efficiency gains will likely encourage more governments to adopt such tools in the coming years. Precisely because adoption is accelerating, this study's central implication is that robust governance laws and frameworks are the precondition for protecting the asylum system, not an afterthought to be added once tools are deployed: a jurisdiction that decides to use AI in asylum procedures should have lifecycle data governance, impact assessment, and published evaluation in place before the first case is touched. Efficiency must not be purchased with fairness, privacy, or procedural safeguards. Given the complexity of asylum claims and the continually changing nature of displacement, AI should augment rather than replace human decision-making in refugee status determination: responsibility for weighing evidence and determining protection needs must remain with authorised human decision-makers, supported by — never substituted by — well-governed AI.

Bibliography

European Commission (2026) European asylum and migration management strategy. Communication to the European Parliament and the Council, 29 January 2026. Brussels: European Commission. Available at: https://home-affairs.ec.europa.eu/european-asylum-and-migration-management-strategy_en

European Parliament (2025) Artificial intelligence in asylum procedures in the EU. Briefing PE 775.861, by C. Dumbrava. Brussels: European Parliamentary Research Service. Available at: [https://www.europarl.europa.eu/thinktank/en/document/EPRS_BRI\(2025\)775861](https://www.europarl.europa.eu/thinktank/en/document/EPRS_BRI(2025)775861)

European Union (2024) Regulation (EU) 2024/1689 of the European Parliament and of the Council laying down harmonised rules on artificial intelligence (Artificial Intelligence Act). Official Journal of the European Union, L series. Available at: <https://eur-lex.europa.eu/eli/reg/2024/1689/oj>

Eurostat (2026) Asylum applications in February 2026. News article, 18 May 2026. Luxembourg: Eurostat. Available at: <https://ec.europa.eu/eurostat/web/products-eurostat-news/w/ddn-20260518-1>

Government of Canada (2024) Guide on the use of generative artificial intelligence. Ottawa: Treasury Board of Canada Secretariat. Available at: <https://www.canada.ca/en/government/system/digital-government/digital-government-innovations/responsible-use-ai/guide-use-generative-ai.html>

Home Office (2025) Evaluation of AI trials in the asylum decision making process. Research and analysis, April 2025. London: Home Office. Available at: <https://www.gov.uk/government/publications/evaluation-of-ai-trials-in-the-asylum-decision-making-process>

Home Office (2026) Facial age estimation: using AI to support initial age decisions (a guide). Guidance, May 2026. London: Home Office. Available at: <https://www.gov.uk/government/publications/facial-age-estimation>

Human Rights Watch (2026) The UK's discriminatory AI experiment on child refugees. 19 June 2026. New York: Human Rights Watch. Available at: <https://www.hrw.org/news/2026/06/19/the-uks-discriminatory-ai-experiment-on-child-refugees>

Information Commissioner's Office (2023) Guidance on AI and data protection. Wilmslow: ICO. Available at: <https://ico.org.uk/for-organisations/uk-gdpr-guidance-and-resources/artificial-intelligence/guidance-on-ai-and-data-protection/>

Innovation, Science and Economic Development Canada (2023) The Artificial Intelligence and Data Act (AIDA) — companion document. Ottawa: ISED. Available at: <https://ised-isde.canada.ca/site/innovation-better-canada/en/artificial-intelligence-and-data-act-aida-companion-document>

Immigration, Refugees and Citizenship Canada (2026) Immigration, Refugees and Citizenship Canada's Artificial Intelligence Strategy. Ottawa: IRCC. Available at: <https://www.canada.ca/en/immigration-refugees-citizenship/corporate/transparency/artificial-intelligence-strategy.html>

National Institute of Standards and Technology (2023) Artificial Intelligence Risk Management Framework (AI RMF 1.0). NIST AI 100-1. Gaithersburg, MD: NIST. Available at: <https://nvlpubs.nist.gov/nistpubs/ai/NIST.AI.100-1.pdf>

Stanford Human-Centered Artificial Intelligence (2026) The AI Index 2026 Annual Report. Stanford, CA: Stanford University. Available at: <https://hai.stanford.edu/ai-index/2026-ai-index-report>

Treasury Board of Canada Secretariat (2023) Directive on Automated Decision-Making. Ottawa: Government of Canada. Available at: <https://www.tbs-sct.canada.ca/pol/doc-eng.aspx?id=32592>

UK Government (2023) A pro-innovation approach to AI regulation. White Paper. London: Department for Science, Innovation and Technology. Available at: <https://www.gov.uk/government/publications/ai-regulation-a-pro-innovation-approach>

UK Government (2025a) Artificial Intelligence Playbook for the UK Government. London: Government Digital Service. Available at: <https://www.gov.uk/government/publications/ai-playbook-for-the-uk-government>

UK Government (2025b) AI Cyber Security Code of Practice. London: Department for Science, Innovation and Technology. Available at: <https://www.gov.uk/government/publications/ai-cyber-security-code-of-practice>

UK Government (2026) Guidelines and best practices for making government datasets ready for AI. London: Government Digital Service and Department for Science, Innovation and Technology. Available at: <https://www.gov.uk/government/publications/making-government-datasets-ready-for-ai/guidelines-and-best-practices-for-making-government-datasets-ready-for-ai>

UNDP (2025) Human Development Report 2025 — A matter of choice: people and possibilities in the age of AI. New York: United Nations Development Programme. Available at: <https://hdr.undp.org/content/human-development-report-2025>

UNHCR (2024) UNHCR's approach to artificial intelligence. Geneva: UNHCR. Available at: <https://www.unhcr.org/digitalstrategy/ai-approach/>

World Bank (2025) AI-powered refugee forecasting: preparing for refugee movements before they happen. Washington, DC: World Bank. Available at: <https://www.worldbank.org/en/topic/fragilityconflictviolence/brief/ai-powered-refugee-forecasting-preparing-for-refugee-movements-before-they-happen>